

Importing the libraries

```
In [1]: import pandas as pd
import numpy as np
import sklearn as sk
from numpy import log,dot,exp,shape
import matplotlib.pyplot as plt
```

Importing the dataset

```
In [2]: dataset=pd.read_csv("diabetes.csv")
```

```
In [3]: dataset
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
...	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

768 rows × 9 columns

Setting the feature and target variable

```
In [4]: x=dataset.iloc[:, :-1].values
y=dataset.iloc[:, -1].values
```

In [5]:

x

Out[5]: array([[6. , 148. , 72. , ..., 33.6 , 0.627, 50.],
[1. , 85. , 66. , ..., 26.6 , 0.351, 31.],
[8. , 183. , 64. , ..., 23.3 , 0.672, 32.],
...,
[5. , 121. , 72. , ..., 26.2 , 0.245, 30.],
[1. , 126. , 60. , ..., 30.1 , 0.349, 47.],
[1. , 93. , 70. , ..., 30.4 , 0.315, 23.]])

In [6]:

y

```
Out[6]: array([[1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1,
0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,
0, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
1, 0,
1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
0, 1,
1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1,
1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0,
0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1,
0, 1,
1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0,
1, 1,
1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0,
0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1,
0, 0,
1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
1, 0,
0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0,
1, 0,
1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0,
1, 0,
0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0,
0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
1, 0,
0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1,
0, 1,
0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0,
1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 0,
1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0,
0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
1, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
1, 0,
0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
```

```

0, 0,
    1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 1,
    0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0,
0, 1,
    0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1,
1, 0,
    0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0,
0, 0,
    0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0,
1, 0,
    1  1  1  0  0  1  1  1  0  1  0  1  0  1  0  0  0  0  1  0

```

Filling the missing data

```

In [7]: from sklearn.impute import SimpleImputer
imputer=SimpleImputer(missing_values=np.nan, strategy='mean')
imputer.fit(x)
x=imputer.transform(x)

```

```
In [8]: x
```

```

Out[8]: array([[ 6.   , 148.   , 72.   , ..., 33.6 , 0.627, 50.   ],
 [ 1.   , 85.   , 66.   , ..., 26.6 , 0.351, 31.   ],
 [ 8.   , 183.   , 64.   , ..., 23.3 , 0.672, 32.   ],
 ...,
 [ 5.   , 121.   , 72.   , ..., 26.2 , 0.245, 30.   ],
 [ 1.   , 126.   , 60.   , ..., 30.1 , 0.349, 47.   ],
 [ 1.   , 93.   , 70.   , ..., 30.4 , 0.315, 23.   ]
])

```

Splitting the dataset into training set and test set(test data set 20%)

```

In [9]: x_train=x[:615,:]
x_test=x[615:,:]
y_train=y[:615]
y_test=y[615:]

```

```
In [10]: x_train
```

```

Out[10]: array([[ 6.   , 148.   , 72.   , ..., 33.6 , 0.627, 50.   ],
 [ 1.   , 85.   , 66.   , ..., 26.6 , 0.351, 31.   ],
 [ 8.   , 183.   , 64.   , ..., 23.3 , 0.672, 32.   ],
 ...,
 [ 7.   , 168.   , 88.   , ..., 38.2 , 0.787, 40.   ],
 [ 6.   , 105.   , 80.   , ..., 32.5 , 0.878, 26.   ],
 [ 11.  , 138.   , 74.   , ..., 36.1 , 0.557, 50.   ]
])

```

```
In [11]: x_test
```

```
Out[11]: array([[ 3.    , 106.    , 72.    , ..., 25.8    , 0.207, 27.    ],
 [ 6.    , 117.    , 96.    , ..., 28.7    , 0.157, 30.    ],
 [ 2.    , 68.     , 62.    , ..., 20.1    , 0.257, 23.    ],
 ...,
 [ 5.    , 121.    , 72.    , ..., 26.2    , 0.245, 30.    ],
 [ 1.    , 126.    , 60.    , ..., 30.1    , 0.349, 47.    ],
 [ 1.    , 93.     , 70.    , ..., 30.4    , 0.315, 23.    ]])
```

In [12]: y_train

```
Out[12]: array([[1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1,
0, 0,
1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0,
0, 1,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0,
1, 0,
1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
0, 0,
1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
0, 1,
1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1,
1, 1,
1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1, 0,
1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0,
0, 1,
0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1,
0, 1,
1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0,
1, 1,
1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0,
0, 0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1,
0, 0,
1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
1, 0,
0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0,
1, 0,
1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0,
1, 0,
0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0,
0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
0, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
1, 0,
0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1,
0, 1,
0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0,
1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 0,
1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0,
0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0,
0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0,
1, 0,
0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0,
```

In [13]: `y_test`

Out[13]: `array([[0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
1, 0,
0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0,
0, 0,
1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1,
0, 0,
1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0,
1, 1,
0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1,
0, 0,
0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0,
0, 1,
0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1,
0])`

Feature Scaling

In [14]: `from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
x_train=sc.fit_transform(x_train)
x_test=sc.transform(x_test)`

In [15]: `x_train`

Out[15]: `array([[0.65218839, 0.84277421, 0.17715826, ..., 0.21380506,
0.43319642, 1.42490668],
[-0.83887428, -1.086291, -0.12931806, ..., -0.66513955,
-0.38742276, -0.1889379],
[1.24861345, 1.9144771, -0.23147683, ..., -1.07949916,
0.56699303, -0.10399871],
...,
[0.95040092, 1.45517586, 0.99442844, ..., 0.79139724,
0.90891769, 0.5755148],
[0.65218839, -0.47388934, 0.58579335, ..., 0.0756852 ,
1.17948416, -0.61363385],
[2.14325106, 0.53657338, 0.27931703, ..., 0.52771386,
0.22506837, 1.42490668]])`

In [16]: `x_test`

Out[16]: `array([[-0.24244921, -0.44326926, 0.17715826, ..., -0.76559036,
-0.81557189, -0.52869466],
[0.65218839, -0.10644835, 1.40306353, ..., -0.40145617,
-0.96423479, -0.27387709],
[-0.54066175, -1.6068324, -0.3336356, ..., -1.48130241,
-0.666909, -0.86845141],
...,
[0.35397585, 0.01603198, 0.17715826, ..., -0.71536496,
-0.70258809, -0.27387709],
[-0.83887428, 0.16913239, -0.43579437, ..., -0.22566724,
-0.39336927, 1.17008912],
[-0.83887428, -0.84133034, 0.07499949, ..., -0.18799819,
-0.49446004, -0.86845141]])`

Logistic Regression Class


```
In [17]: #defining the logistic regression class
class LogisticRegression:

    #sigmoid function used for calculating the hypothesis
    def sigmoid(self,z):
        sig = 1/(1+exp(-z))
        return sig

    def initialize(self,X):
        weights = np.zeros((shape(X)[1]+1,1))
        X = np.c_[np.ones((shape(X)[0],1)),X]
        return weights,X


    #training the model
    def fit(self,X,y,alpha=0.001,iter=400):
        weights,X = self.initialize(X)
        def cost(theta):
            z = dot(X,theta)
            cost0 = y.T.dot(log(self.sigmoid(z)))
            cost1 = (1-y).T.dot(log(1-self.sigmoid(z)))
            cost = -((cost1 + cost0))/len(y)
            return cost
        cost_list = np.zeros(iter,)
        for i in range(iter):
            weights = weights - alpha*dot(X.T,self.sigmoid(dot(X,
            cost_list[i] = cost(weights)
        self.weights = weights
        return cost_list

    #for result prediction
    def predict(self,X):
        z = dot(self.initialize(X)[1],self.weights)
        lis = []
        for i in self.sigmoid(z):
            if i>0.5:
                lis.append(1)
            else:
                lis.append(0)
        return lis
```



$$\text{cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

if $y = 1$ if $y = 0$



Model Training

```
In [18]: #object creation of the LogisticRegression class  
classifier= LogisticRegression()
```

```
In [19]: #fitting our model on training set
model= classifier.fit(x_train,y_train)
```

Model Prediction

```
In [20]: #predicting the rersults on test set
y_pred = classifier.predict(x_test)
#y_pred here is a list
```

```
In [21]: #converting the pandas series to a list
         y_test=y_test.tolist()
```

```
In [22]: y_pred
```

[illegible]

```
Out[23]: [0,  
          0,  
          0,  
          1,  
          1,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          0,  
          1,  
          0,  
          0,  
          0,  
          0,
```

A scatter plot titled "y_pred vs y_test" showing the relationship between predicted values (y_pred) on the x-axis and test values (y_test) on the y-axis. The x-axis ranges from 0.0 to 1.0, and the y-axis ranges from 0.0 to 1.0. There are four data points plotted as red dots: (0.0, 0.0), (0.0, 1.0), (1.0, 0.0), and (1.0, 1.0).

Defining F1_score

```
In [25]: def F1_score(y,y_pred):
          tp,tn,fp,fn = 0,0,0,0
          for i in range(len(y)):
              if y[i] == 1 and y_pred[i] == 1:
                  tp += 1
              elif y[i] == 1 and y_pred[i] == 0:
                  fp += 1
              elif y[i] == 0 and y_pred[i] == 1:
                  fn += 1
              elif y[i] == 0 and y_pred[i] == 0:
                  tn += 1
          precision = tp/(tp+fp)
          recall = tp/(tp+fn)
          f1_score = 2*precision*recall/(precision+recall)
          return f1_score
```

F1 is calculated as follows:

$$F_1 = 2 * \frac{precision * recall}{precision + recall}$$

where:

$$precision = \frac{TP}{TP + FP}$$

$$recall = \frac{TP}{TP + FN}$$

In "macro" F1 a separate F1 score is calculated for each `species` value and then averaged.

F1_Score Metric value

```
In [26]: #calculating the f1_score for checking the accuracy
          f1_score=F1_score(y_test,y_pred)
```

```
In [27]: f1_score
```

```
Out[27]: 0.6236559139784946
```

Defining The Accuracy

```
In [28]: def accuracy(y,y_pred):
          tp,tn,fp,fn=0,0,0,0
          for i in range(len(y)):
              if y[i]==1 and y_pred[i]==1:
                  tp+=1
              elif y[i]==1 and y_pred[i]==0:
                  fp+=1
              elif y[i]==0 and y_pred[i]==1:
                  fn+=1
              elif y[i]==0 and y_pred[i]==0:
                  tn+=1
          accuracy_score=(tp+tn)/(tp+tn+fp+fn)
          return accuracy_score
```

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Accuracy Metric Value

```
In [29]: #calculating the accuracy
          accuracy=accuracy(y_test,y_pred)
          accuracy
```

Out[29]: 0.7712418300653595